

RAVITEJA THOLUPUNOORI

Software Engineer III | Sr. Data Engineer

raviteja0096@gmail.com | +91-9542710096 | [linkedin.com/in/traviteja](https://www.linkedin.com/in/traviteja) | github.com/traviteja-git | traviteja.com

PROFESSIONAL SUMMARY

9+ years of experience in data and AI platform engineering, specialising in building production-grade data pipelines, LLM-based platforms, and multi-cloud infrastructure. Delivered a brand matching engine from zero to production in a single sprint, an LLM-based campaign generation platform, and vector embedding pipelines — as part of a 90+ repo multi-product SaaS platform on GCP, Azure, and Databricks.

Architected the core Identity Resolution pipeline processing 800 million records with 50% reduction in processing time, modernised legacy .NET to Apache Spark (70% performance gain), and extended the platform to AWS and Azure within 3 months. Apache Airflow Champion (Astronomer, Cohort 3).

TECHNICAL SKILLS

AI / Data	Claude, Copilot, ChatGPT, GenAI Prompt Engineering, Apache Spark, Apache Airflow, Databricks, BigQuery
Languages	Python, C#, SQL
Cloud	GCP, Azure, AWS
Containers	Docker, Kubernetes (GKE / AKS), KEDA, Argo CD, Helm
API & Backend	FastAPI, RESTful APIs
Databases	PostgreSQL, MongoDB, SQL Server
Prior Stack	ASP.NET MVC, .NET Core, SSIS, Talend

CERTIFICATIONS

- Professional Cloud Architect · Professional Developer · Associate Cloud Engineer · Cloud Leader | Google Cloud Platform
- Apache Airflow Fundamentals · Apache Airflow Champion (Cohort 3) | Astronomer
- SI Architect · SI Associate | MongoDB
- Microsoft Certified Developer — Programming in C# | Microsoft
- AWS Technical Professional | Amazon Web Services
- Certified GenAI Engineer | Deloitte

EXPERIENCE

Software Engineer III — Deloitte Consulting India PVT. Ltd. Aug 2024 – Present | Hyderabad

ConvergeCONSUMER — Customer & Marketing / Innovation & Technology

Customer Growth Opportunity Platform · Python · FastAPI · PostgreSQL · Kubernetes · GCP · Databricks · KEDA · ArgoCD

- Stabilized the Customer Growth core module by resolving 20+ functional and platform issues, enabling successful end-to-end validation and production readiness across development, demo, and staging environments.
- Designed the HLD and led containerisation of the Customer Growth app — converting a manual workflow into a fully automated, containerised deployment via GitHub Actions CI/CD.
- Piloted KEDA HTTP-trigger autoscaling for Customer Growth, introducing an event-driven scaling pattern that reduced operational complexity compared to the previous DB-trigger approach.
- Delivered **semantic search for opportunity matching** by implementing a Gen Model-based search layer in the Customer Growth API.
- Led a 5-member hackathon team to accelerate cc-utilities adoption across products, completing the majority of migration activities and establishing a roadmap for product teams to finish remaining work.
- Automated publishing of DB Connector and GenAI artifacts using Twine-based workflows, streamlining build and artifact deployment pipelines.
- Guided the Subscriber Activation team on Campaigns API delivery — clarified technical dependencies, ownership boundaries, and resolution paths for cross-team execution.

★ *Applause Award — delivering MVP for a client within the timeline*

Customer Segmentation · Python · PySpark · GCP · BigQuery · Kubernetes · Pub/Sub

- Led the integration of Model Ready Data (MRD) into the Customer Segmentation platform, eliminating manual Jupyter-based preprocessing and enabling fully automated end-to-end ML workflows.
- Designed a pub/sub-based architecture to support multi-application reuse — enabling scalable, decoupled data processing across teams without manual intervention.
- Designed and implemented a PySpark-based Data Prep Pipeline with advanced data quality checks; delivered MRD pipeline integration, dependency fixes, and UI improvements end-to-end through UAT.
- Increased automated test coverage to approximately 90% across segmentation and clustering modules, improving release confidence and maintainability.

- Refactored the platform to use cc-utilities shared library, implemented RBAC action changes, and built the file manager feature for multi-user access governance.

Brand Identity · Python · Big Query · GCP · Docker · CASS

- Built and delivered an end-to-end Brand Identity matching MVP within a single sprint, including matching logic, CASS integration, Big Query data layer, testing, containerization, and CI/CD.
- CASS integration standardises client addresses before matching — this was the main accuracy gap in the previous approach.

Product Engineer 2 — Deloitte Consulting India Pvt. Ltd.

Jul 2021 – Jul 2024 | Hyderabad

ConvergeCONSUMER — Customer & Marketing / Innovation & Technology

Identity Resolution Platform · Apache Spark · Airflow · Python · C# · GCP · Azure · AWS · Databricks · Kubernetes · MongoDB

- Migrated Identity Resolution infrastructure to support AWS and Azure within **3 months**, enabling automated single-click deployment across three cloud providers.
- Modernised data processing from legacy .NET to Apache Spark — **70% increase in job performance**; designed and implemented the core IDR pipeline processing **800 million records** with **50% reduction in processing time** vs. the existing solution.
- Built an Airflow-orchestrated batch ingestion pipeline for the microservice suite (Cleanse, Convert, CASS, Pinning, Stitch) — **25% cost reduction, 20% scalability increase** — covering data hygiene, address normalisation, identity resolution, and data governance.
- Extended the pipeline with Databricks delta table orchestration — per-run statistics tracking, master record count validation, and Azure + Databricks filesystem integration via cc-utilities for true multi-cloud IDR execution.
- Designed and built an IDR Databricks Dashboard within the Databricks environment — eliminating the need for a separate UI, reducing infrastructure cost, and simplifying overall architecture.
- Productized all microservices as both console applications and REST APIs, deployed on Azure and GCP Kubernetes.
- Extended the pipeline with Databricks delta table orchestration — per-run statistics tracking, master record count validation, and Azure + Databricks filesystem integration via cc-utilities for true multi-cloud IDR execution.
- Maintained and operated **20+ Apache Airflow DAGs** on GCP Cloud Composer — custom Docker image builds, Kubernetes pod configs, pre-commit enforcement, and pipeline stability fixes.

★ *Applause Award* — .NET to Spark migration delivered within the agreed timeframe

★ *Outstanding Performance Award* — end-to-end delivery and ongoing production support

Full Stack Developer — Strategic Technology Resources International PVT. Ltd. Apr 2020 – Jun 2021 | Hyderabad

- Built modular single-page UIs with AngularJS and Node.js for IntellispeX; acted as client-facing interface for requirements and backend API integration.
- Developed a low-code ATS on Zoho Creator — job posting, applicant management, and keyword candidate search across social platforms.
- Built and migrated configurable .NET/Java automation apps to Azure for SPS Commerce; developed complex SQL procedures fetching from CSV and XML sources.

Business Technology Analyst — Deloitte Consulting India PVT. Ltd.

Apr 2018 – Mar 2020 | Hyderabad

LGE CDP Implementation · GCP · PostgreSQL · Talend · Apache Airflow · Python · Docker

- Built a configurable Talend automation job processing >400GB dovetail files (GCP download → decrypt → unzip → convert → cleanse → reload) — **18–20 day manual process reduced to 2–3 days**; optimised a critical data feed from **>13 hours to 1 hour**.

★ *Applause Award* — took ownership across multiple tech stacks, delivered on time

Tech Modernisation · C# · .NET Core · MongoDB · Docker

- Built real-time microservices (Convert, CASS, Pinning, Data Loader, Drops) for identity resolution and data hygiene. **Data Loader:** 100K records in 1.3 mins (was 30 mins). **Drops:** 1M records in 2 mins (was >60 hours).

★ *Spot Award* — exceptional application performance improvements

Software Engineer — Capgemini Technology Services India Limited

Oct 2016 – Apr 2018 | Bangalore

- Developed and maintained the RBC (Royal Bank of Canada) Sales Platform — a web-based product sales system on C# / .NET MVC.
- Built configurable environment-identification files for multi-environment deployment and testing across dev, staging, and production.

EDUCATION

B.Tech, Electronics & Communication Engineering — CVR College of Engineering, JNTUH (81%)

2016